

Introduction

In this project, you will build a neural network of your own design to evaluate the CIFAR-10 dataset. Our target accuracy is 70%, but any accuracy over 50% is a great start. Some of the benchmark results on CIFAR-10 include:

78.9% Accuracy | [Deep Belief Networks; Krizhevsky, 2010](#)

90.6% Accuracy | [Maxout Networks; Goodfellow et al., 2013](#)

96.0% Accuracy | [Wide Residual Networks; Zagoruyko et al., 2016](#)

99.0% Accuracy | [GPipe; Huang et al., 2018](#)

98.5% Accuracy | [Rethinking Recurrent Neural Networks and other Improvements for ImageClassification; Nguyen et al., 2020](#)

Research with this dataset is ongoing. Notably, many of these networks are quite large and quite expensive to train.

Imports

```
In [1]: ## This cell contains the essential imports you will need - DO NOT CHANGE THE CONTENTS! ##
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
import numpy as np
```

Load the Dataset

Specify your transforms as a list first. The transforms module is already loaded as `transforms`.

CIFAR-10 is fortunately included in the torchvision module. Then, you can create your dataset using the `CIFAR10` object from `torchvision.datasets` ([the documentation is available here](#)). Make sure to specify `download=True` !

Once your dataset is created, you'll also need to define a `DataLoader` from the `torch.utils.data` module for both the train and the test set.

```
In [2]: from torch.utils.data import DataLoader
from torchvision.datasets import CIFAR10
```

```
In [3]: dataset_path = "C:\\Users\\Vivek\\Udacity\\DeepLearningUsingPyTorch\\Data\\"
batch_size = 32

torch.manual_seed(42)
np.random.seed(42)

# From https://pytorch.org/vision/stable/models.html
mean = [0.485, 0.456, 0.406]
std = [0.229, 0.224, 0.225]

# Define transforms
# Note to self: Try the RandomCrop or CenterCrop transforms to improve accuracy if needed
transforms_train = transforms.Compose([
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(20),
    transforms.Resize(224), # Min size for pretrained
    transforms.ToTensor(),
    transforms.Normalize(mean=mean, std=std)
])
transforms_test = transforms.Compose([
    transforms.Resize(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=mean, std=std)
])

# Create training set and define training dataloader
trainset = CIFAR10(root=dataset_path, train=True, transform=transforms_train, download=True)
trainloader = DataLoader(dataset=trainset, batch_size=batch_size, shuffle=True)

# Create test set and define test dataloader
testset = CIFAR10(root=dataset_path, train=False, transform=transforms_test, download=True)
testloader = DataLoader(dataset=testset, batch_size=batch_size)

# The 10 classes in the dataset
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Files already downloaded and verified
Files already downloaded and verified

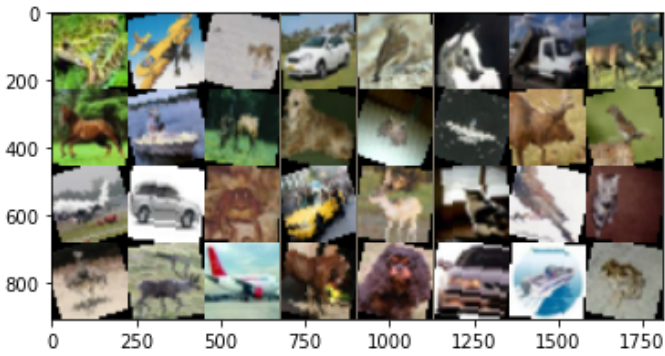
```
In [4]: # Source: https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html

import matplotlib.pyplot as plt
import numpy as np
import os

# functions to show an image
denormalize = transforms.Compose([
    transforms.Normalize(mean=[-m/s for (m, s) in zip(mean, std)], std=[1.0/s for s in std])
])
```

```
def imshow(img):
    img = denormalize(img)
    npimg = img.numpy()
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
    plt.show()
```

```
In [5]: images, labels = next(iter(trainloader))
imshow(torchvision.utils.make_grid(images))
print(''.join(f"{str(os.linesep) + str(j//8) + ': ' if j%8 == 0 else ', '}{classes[labels[j]]:5s}" for j in range(batch_size)))
```



0: frog , plane, deer , car , bird , horse, truck, deer
1: horse, ship , deer , dog , frog , plane, deer , bird
2: plane, car , frog , car , deer , cat , bird , cat
3: bird , deer , plane, horse, dog , car , ship , frog

Explore the Dataset

Using matplotlib, numpy, and torch, explore the dimensions of your data.

You can view images using the `show5` function defined below – it takes a data loader as an argument. Remember that normalized images will look really weird to you! You may want to try changing your transforms to view images. Typically using no transforms other than `toTensor()` works well for viewing – but not as well for training your network. If `show5` doesn't work, go back and check your code for creating your data loaders and your training/test sets.

```
In [6]: def show5(img_loader):
        dataiter = iter(img_loader)

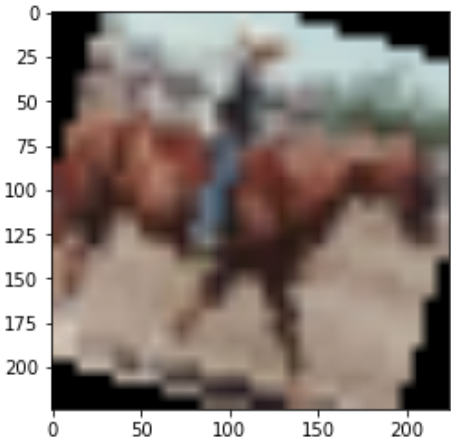
        batch = next(dataiter)
        labels = batch[1][0:5]
        images = batch[0][0:5]
        for i in range(5):
            print(classes[labels[i]])

            image = denormalize(images[i])
            npimg = image.numpy()
            plt.imshow(np.transpose(npimg, (1, 2, 0)))
            plt.show()
```

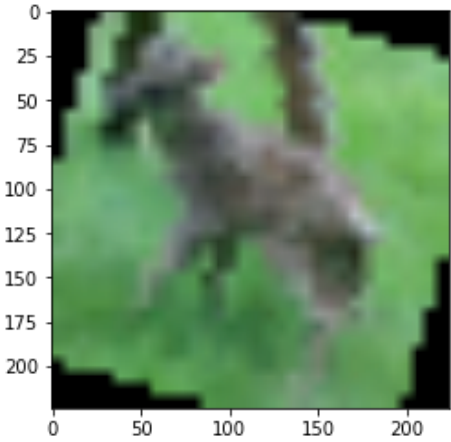
```
In [7]: # Explore data
print("TRAIN DATA")
show5(trainloader)
print("TEST DATA")
show5(testloader)
```

TRAIN DATA

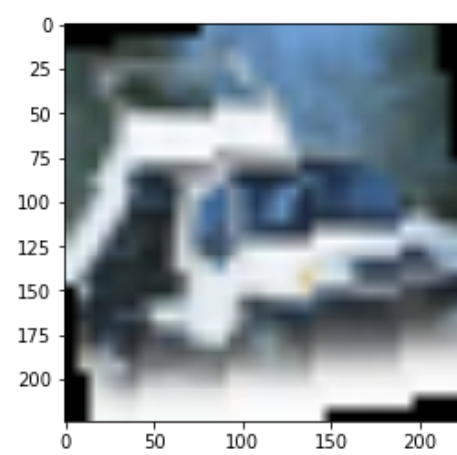
horse



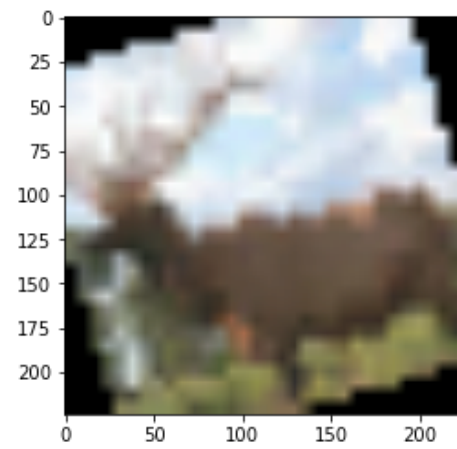
deer



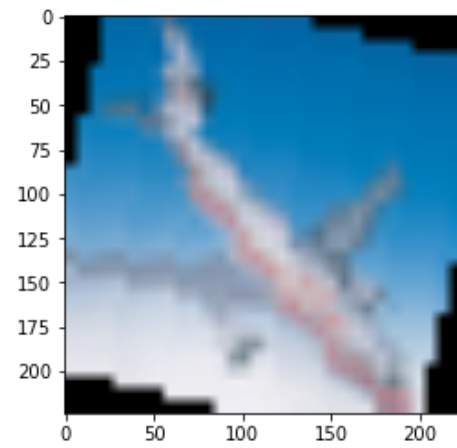
truck



deer

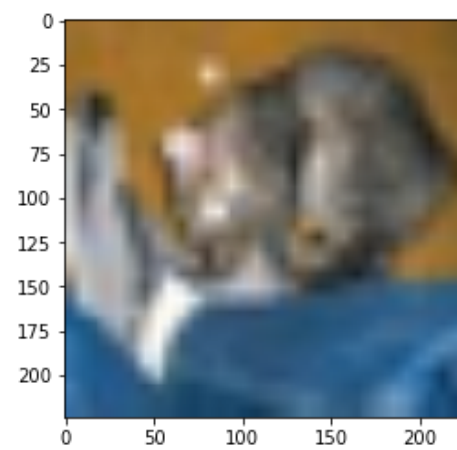


plane




TEST DATA

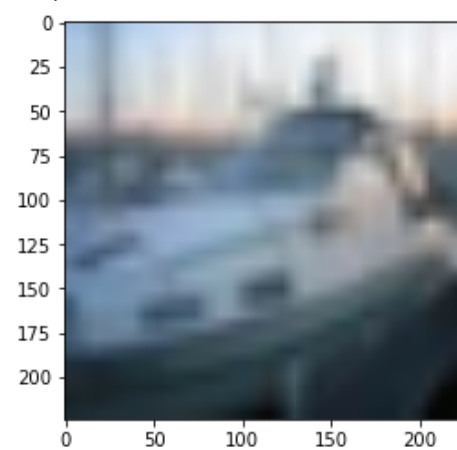
cat



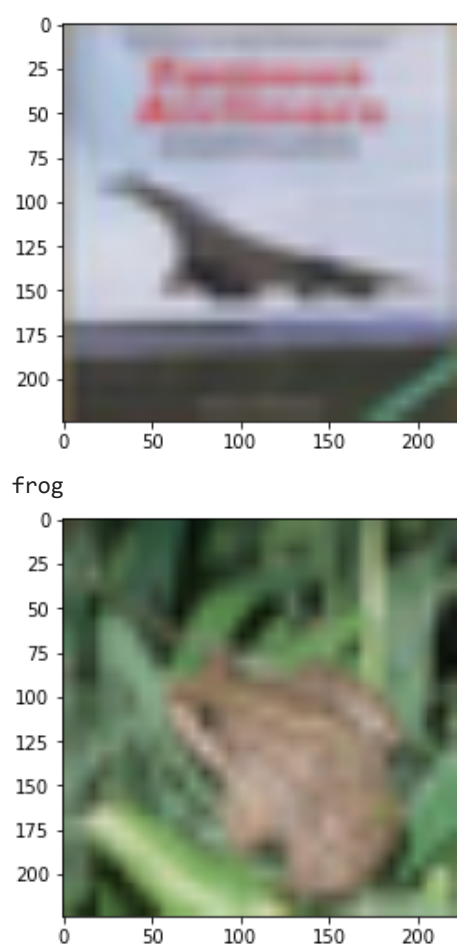
ship



ship



plane



Build your Neural Network

Using the layers in `torch.nn` (which has been imported as `nn`) and the `torch.nn.functional` module (imported as `F`), construct a neural network based on the parameters of the dataset. Feel free to construct a model of any architecture – feedforward, convolutional, or even something more advanced!

In [8]: *# NOTE: I also tried to train my own model but the accuracy was not great (51.37%).
Code and results for that model are after the *Make a recommendation* section*

```
model = torchvision.models.resnet152(pretrained=True)

for param in model.parameters():
    param.requires_grad = False

classifier = nn.Sequential(
    nn.Linear(2048, 1024),
    nn.ReLU(),
    nn.Dropout(p=0.2),
    nn.Linear(1024, 512),
    nn.ReLU(),
    nn.Dropout(p=0.2),
    nn.Linear(512, 10),
    nn.LogSoftmax(dim=1)
)

model.fc = classifier
```

Specify a loss function and an optimizer, and instantiate the model.

If you use a less common loss function, please note why you chose that loss function in a comment.

In [9]:

```
from torch.optim import Adam
from torch.nn import NLLLoss

optimizer = Adam(params=model.fc.parameters(), lr=0.001)
criterion = NLLLoss()
```

Running your Neural Network

Use whatever method you like to train your neural network, and ensure you record the average loss at each epoch. Don't forget to use `torch.device()` and the `.to()` method for both your model and your data if you are using GPU!

If you want to print your loss during each epoch, you can use the `enumerate` function and print the loss after a set number of batches. 250 batches works well for most people!

In [10]:

```
device = "cuda:0" if torch.cuda.is_available() else "cpu"
model.to(device)
print(f"Using device '{device}')
```

Using device 'cuda:0'

In [11]:

```
EPOCHS = range(20)

per_epoch_train_loss = []
per_epoch_test_loss = []
per_epoch_accuracy = []

for epoch in EPOCHS:
    total_loss = 0

    for batch, (images, labels) in enumerate(trainloader):
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
```

```

log_ps = model.forward(images)
loss = criterion(log_ps, labels)
loss.backward()
optimizer.step()

total_loss += loss

# if batch%250 == 0:
#     print(f"Epoch {epoch:2} Batch {batch:4} Loss: {loss:.3f} Avg Loss: {total_loss/float(batch+1):.3f}")

per_epoch_train_loss.append(total_loss.item()/float(batch+1)) # Per batch Loss

# Compute Test Loss
with torch.no_grad():
    model.eval()

total_loss = 0
total_correct = 0
count = 0

for batch, (images, labels) in enumerate(testloader):
    images, labels = images.to(device), labels.to(device)
    log_ps = model.forward(images)
    total_loss += criterion(log_ps, labels)
    ps, pred_label = log_ps.topk(1, dim=1)

    total_correct += torch.sum(pred_label.view(labels.shape) == labels)
    count += len(labels)

per_epoch_test_loss.append(total_loss.item() / float(batch))
per_epoch_accuracy.append(100. * total_correct.item() / float(count))
print(f"Epoch {epoch:2} Train Loss: {per_epoch_train_loss[-1]:.3f} Test Loss: {per_epoch_test_loss[-1]:.3f} Accuracy {per_epoch_accuracy[-1]:.3f}")

model.train()

# Checkpoint at each epoch
checkpoint_path = f"C:\\Users\\Vivek\\Udacity\\DeepLearningUsingPyTorch\\CIFAR10_classifier_checkpoints\\CIFAR10_classifier_model_epoch{epoch}.pt"
torch.save(model.state_dict(), checkpoint_path)

```

```

Epoch 0 Train Loss: 1.178 Test Loss: 0.727 Accuracy 75.10%
Epoch 1 Train Loss: 1.020 Test Loss: 0.656 Accuracy 77.35%
Epoch 2 Train Loss: 0.981 Test Loss: 0.678 Accuracy 76.58%
Epoch 3 Train Loss: 0.961 Test Loss: 0.627 Accuracy 78.10%
Epoch 4 Train Loss: 0.949 Test Loss: 0.610 Accuracy 79.16%
Epoch 5 Train Loss: 0.922 Test Loss: 0.583 Accuracy 79.80%
Epoch 6 Train Loss: 0.923 Test Loss: 0.612 Accuracy 79.39%
Epoch 7 Train Loss: 0.911 Test Loss: 0.589 Accuracy 79.80%
Epoch 8 Train Loss: 0.904 Test Loss: 0.606 Accuracy 79.20%
Epoch 9 Train Loss: 0.901 Test Loss: 0.580 Accuracy 80.39%
Epoch 10 Train Loss: 0.886 Test Loss: 0.578 Accuracy 80.00%
Epoch 11 Train Loss: 0.884 Test Loss: 0.573 Accuracy 79.96%
Epoch 12 Train Loss: 0.878 Test Loss: 0.569 Accuracy 79.98%
Epoch 13 Train Loss: 0.866 Test Loss: 0.551 Accuracy 80.85%
Epoch 14 Train Loss: 0.879 Test Loss: 0.567 Accuracy 80.49%
Epoch 15 Train Loss: 0.864 Test Loss: 0.535 Accuracy 81.68%
Epoch 16 Train Loss: 0.862 Test Loss: 0.541 Accuracy 81.29%
Epoch 17 Train Loss: 0.852 Test Loss: 0.550 Accuracy 80.87%
Epoch 18 Train Loss: 0.851 Test Loss: 0.543 Accuracy 81.43%
Epoch 19 Train Loss: 0.852 Test Loss: 0.543 Accuracy 81.13%


```

Plot the training loss (and validation loss/accuracy, if recorded).

```

In [12]: plt.plot(per_epoch_test_loss)
plt.plot(per_epoch_train_loss)
plt.plot([x/100. for x in per_epoch_accuracy])
plt.legend(["Test Loss", "Train Loss", "Accuracy"])
plt.show()

```



Testing your model

Using the previously created `DataLoader` for the test set, compute the percentage of correct predictions using the highest probability prediction.

If your accuracy is over 70%, great work! This is a hard task to exceed 70% on.

If your accuracy is under 45%, you'll need to make improvements. Go back and check your model architecture, loss function, and optimizer to make sure they're appropriate for an image classification task.

```

In [13]: # Already did this in the model training.

```

Saving your model

Using `torch.save`, save your model for future loading.

```
In [14]: model_path = "C:\\Users\\Vivek\\Udacity\\DeepLearningUsingPyTorch\\CIFAR10_classifier_model.pth"
torch.save(model.state_dict(), model_path)
```

Make a Recommendation

Based on your evaluation, what is your recommendation on whether to build or buy? Explain your reasoning below.

```
In [15]: torch.cuda.get_device_properties(0)
```

```
Out[15]: _CudaDeviceProperties(name='NVIDIA GeForce GTX 1660 SUPER', major=7, minor=5, total_memory=6143MB, multi_processor_count=22)
```

Recommendation: BUILD

Since this is an image classification problem with typical (commonly seen) classes, there already are several publicly available pre-trained models as well as public datasets. By utilizing the pretrained feature detection portion of these models, we can quickly replace and retrain just the classifier part of the network to get more than 80% accuracy (81.68% on epoch 15). Given that I was able to throw this together in an afternoon with just a 1660 SUPER, we can certainly get better accuracy by looking at the latest research and experimenting with pre-trained image based models other than the RESNET152 that was used here.

A few notes:

- Using better hardware such as a local machine with 4 or 8 NVidia 3080 or a cloud setup would help further in iterating more quickly.
- The pre-trained model requires an input of minimum size 224 x 224 which is much bigger than our images (32 x 32), so updating the model input shapes and training from scratch might actually not be too slow.
- Creating a representative subset of the dataset would help speed up experimentation
- The graph indicates that both test and train loss are still dropping so it might be interesting to train longer until the model starts overfitting. However, the rate of improvement has dropped significantly (as expected) so it may not be worth spending more time on it.

Pre-trained vs self-built model

Using a pretrained model was very useful in that my GPU only has 6GB of memory and it is very likely that all the weights/biases and gradients of such a large model will probably not fit. Since I'm only retraining a small part that is the classifier, I was able to do it much more quickly.

I also tried training my own model (see below) and was able to get a not that great accuracy of 61.58%. I would probably need to spend some more time improving this model to get anywhere close to the accuracy of the pre-trained network.

Interestingly, the smaller model had higher GPU, CPU as well as disk utilization. While the larger model was around 90% GPU (due to occasional peaks and valleys) and 40% CPU and lower disk use, the smaller model pegged the GPU at 100% with the CPU at 65% and disk at 75%. Perhaps I should have added timing per step to see if the smaller model trains faster and therefore keeps the CPU and disk (dataloader) busy. It is also possible that since I'm retraining only a small part of the larger model, it doesn't have enough compute to keep the GPU at 100%. Might be a good idea to run the NVidia performance analysis tool to validate one of these hypothesis, but any notes from the reviewer would be welcome :)

Attempt at training my own model (not the pretrained one)

```
In [24]: class CIFARClassifier(nn.Module):
def __init__(self):
    super(CIFARClassifier, self).__init__()
    self.conv1 = nn.Conv2d(in_channels=3, out_channels=128, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3))
    self.bn1 = nn.BatchNorm2d(128)
    self.conv2 = nn.Conv2d(in_channels=128, out_channels=64, kernel_size=(5, 5), stride=(2, 2), padding=(2, 2))
    self.bn2 = nn.BatchNorm2d(64)
    self.conv3 = nn.Conv2d(in_channels=64, out_channels=32, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1))
    self.bn3 = nn.BatchNorm2d(32)
    self.conv4 = nn.Conv2d(in_channels=32, out_channels=8, kernel_size=(1, 1), stride=(2, 2), padding=(1, 1))
    self.bn4 = nn.BatchNorm2d(8)

    self.classifier = nn.Sequential(
        nn.Linear(1800, 450),
        nn.ReLU(),
        nn.Dropout(p=0.2),
        nn.Linear(450, 10),
        nn.LogSoftmax(dim=1)
    )

def forward(self, images):
    out = self.conv1(images)
    out = self.bn1(out)
    out = self.conv2(out)
    out = self.bn2(out)
    out = self.conv3(out)
    out = self.bn3(out)
    out = self.conv4(out)
    out = self.bn4(out)
    out = F.relu(out)
    out = F.dropout(out, p=0.2)
    out = self.classifier(out.view(out.shape[0], -1))
    return out

model = CIFARClassifier()
```

Specify a loss function and an optimizer, and instantiate the model.

If you use a less common loss function, please note why you chose that loss function in a comment.

```
In [25]: from torch.optim import Adam
        from torch.nn import NLLLoss

optimizer = Adam(params=model.parameters(), lr=0.001)
# scheduler = MultiStepLR(optimizer, milestones=[1,8], gamma=0.1)
criterion = NLLLoss()
```

Running your Neural Network

Use whatever method you like to train your neural network, and ensure you record the average loss at each epoch. Don't forget to use `torch.device()` and the `.to()` method for both your model and your data if you are using GPU!

If you want to print your loss during each epoch, you can use the `enumerate` function and print the loss after a set number of batches. 250 batches works well for most people!

```
In [26]: device = "cuda:0" if torch.cuda.is_available() else "cpu"
        model.to(device)
        print(f"Using device '{device}'")
```

Using device 'cuda:0'

```
In [27]: EPOCHS = range(20)

per_epoch_train_loss = []
per_epoch_test_loss = []
per_epoch_accuracy = []

for epoch in EPOCHS:
    total_loss = 0

    for batch, (images, labels) in enumerate(trainloader):
        images, labels = images.to(device), labels.to(device)
        optimizer.zero_grad()
        log_ps = model.forward(images)
        loss = criterion(log_ps, labels)
        loss.backward()
        optimizer.step()

        total_loss += loss

        if batch%250 == 0:
            print(f"Epoch {epoch:2} Batch {batch:4} Loss: {loss:.3f} Avg Loss: {total_loss/float(batch+1):.3f}")

    per_epoch_train_loss.append(total_loss.item()/float(batch+1)) # Per batch Loss

    # Compute Test Loss
    with torch.no_grad():
        model.eval()

        total_loss = 0
        total_correct = 0
        count = 0

        for batch, (images, labels) in enumerate(testloader):
            images, labels = images.to(device), labels.to(device)
            log_ps = model.forward(images)
            total_loss += criterion(log_ps, labels)
            ps, pred_label = log_ps.topk(1, dim=1)

            total_correct += torch.sum(pred_label.view(labels.shape) == labels)
            count += len(labels)

        per_epoch_test_loss.append(total_loss.item() / float(batch))
        per_epoch_accuracy.append(100. * total_correct.item() / float(count))
        print(f"Epoch {epoch:2} Train Loss: {per_epoch_train_loss[-1]:.3f} Test Loss: {per_epoch_test_loss[-1]:.3f} Accuracy {per_epoch_accuracy[-1]:.3f}")

    model.train()
```

Epoch	0	Batch	0	Loss:	2.320	Avg	Loss:	2.320
Epoch	0	Batch	250	Loss:	1.611	Avg	Loss:	1.859
Epoch	0	Batch	500	Loss:	1.593	Avg	Loss:	1.782
Epoch	0	Batch	750	Loss:	1.358	Avg	Loss:	1.726
Epoch	0	Batch	1000	Loss:	1.428	Avg	Loss:	1.691
Epoch	0	Batch	1250	Loss:	1.261	Avg	Loss:	1.663
Epoch	0	Batch	1500	Loss:	1.364	Avg	Loss:	1.636
Epoch	0	Train		Loss:	1.633	Test	Loss:	1.521 Accuracy 46.26%
Epoch	1	Batch	0	Loss:	1.207	Avg	Loss:	1.207
Epoch	1	Batch	250	Loss:	1.802	Avg	Loss:	1.480
Epoch	1	Batch	500	Loss:	1.194	Avg	Loss:	1.458
Epoch	1	Batch	750	Loss:	1.391	Avg	Loss:	1.454
Epoch	1	Batch	1000	Loss:	1.444	Avg	Loss:	1.446
Epoch	1	Batch	1250	Loss:	1.367	Avg	Loss:	1.441
Epoch	1	Batch	1500	Loss:	1.212	Avg	Loss:	1.436
Epoch	1	Train		Loss:	1.436	Test	Loss:	1.337 Accuracy 52.93%
Epoch	2	Batch	0	Loss:	1.276	Avg	Loss:	1.276
Epoch	2	Batch	250	Loss:	1.515	Avg	Loss:	1.379
Epoch	2	Batch	500	Loss:	1.616	Avg	Loss:	1.382
Epoch	2	Batch	750	Loss:	1.361	Avg	Loss:	1.379
Epoch	2	Batch	1000	Loss:	1.399	Avg	Loss:	1.378
Epoch	2	Batch	1250	Loss:	1.337	Avg	Loss:	1.376
Epoch	2	Batch	1500	Loss:	1.067	Avg	Loss:	1.375
Epoch	2	Train		Loss:	1.376	Test	Loss:	1.286 Accuracy 54.25%
Epoch	3	Batch	0	Loss:	1.788	Avg	Loss:	1.788
Epoch	3	Batch	250	Loss:	1.273	Avg	Loss:	1.367
Epoch	3	Batch	500	Loss:	1.492	Avg	Loss:	1.355
Epoch	3	Batch	750	Loss:	1.203	Avg	Loss:	1.348
Epoch	3	Batch	1000	Loss:	1.445	Avg	Loss:	1.351
Epoch	3	Batch	1250	Loss:	1.273	Avg	Loss:	1.343
Epoch	3	Batch	1500	Loss:	0.860	Avg	Loss:	1.340
Epoch	3	Train		Loss:	1.341	Test	Loss:	1.251 Accuracy 56.05%
Epoch	4	Batch	0	Loss:	1.761	Avg	Loss:	1.761
Epoch	4	Batch	250	Loss:	1.428	Avg	Loss:	1.319
Epoch	4	Batch	500	Loss:	0.876	Avg	Loss:	1.306
Epoch	4	Batch	750	Loss:	1.516	Avg	Loss:	1.311
Epoch	4	Batch	1000	Loss:	1.045	Avg	Loss:	1.311
Epoch	4	Batch	1250	Loss:	0.980	Avg	Loss:	1.312
Epoch	4	Batch	1500	Loss:	1.175	Avg	Loss:	1.309
Epoch	4	Train		Loss:	1.309	Test	Loss:	1.232 Accuracy 55.70%
Epoch	5	Batch	0	Loss:	1.227	Avg	Loss:	1.227
Epoch	5	Batch	250	Loss:	1.079	Avg	Loss:	1.278
Epoch	5	Batch	500	Loss:	1.095	Avg	Loss:	1.285
Epoch	5	Batch	750	Loss:	1.392	Avg	Loss:	1.291
Epoch	5	Batch	1000	Loss:	1.231	Avg	Loss:	1.296
Epoch	5	Batch	1250	Loss:	1.462	Avg	Loss:	1.296
Epoch	5	Batch	1500	Loss:	1.140	Avg	Loss:	1.294
Epoch	5	Train		Loss:	1.294	Test	Loss:	1.235 Accuracy 56.20%
Epoch	6	Batch	0	Loss:	1.314	Avg	Loss:	1.314
Epoch	6	Batch	250	Loss:	1.250	Avg	Loss:	1.262
Epoch	6	Batch	500	Loss:	1.154	Avg	Loss:	1.263
Epoch	6	Batch	750	Loss:	1.119	Avg	Loss:	1.267
Epoch	6	Batch	1000	Loss:	1.740	Avg	Loss:	1.271
Epoch	6	Batch	1250	Loss:	1.389	Avg	Loss:	1.267
Epoch	6	Batch	1500	Loss:	0.991	Avg	Loss:	1.267
Epoch	6	Train		Loss:	1.267	Test	Loss:	1.228 Accuracy 56.04%
Epoch	7	Batch	0	Loss:	1.166	Avg	Loss:	1.166
Epoch	7	Batch	250	Loss:	1.143	Avg	Loss:	1.265
Epoch	7	Batch	500	Loss:	1.411	Avg	Loss:	1.252
Epoch	7	Batch	750	Loss:	1.381	Avg	Loss:	1.251
Epoch	7	Batch	1000	Loss:	1.189	Avg	Loss:	1.252
Epoch	7	Batch	1250	Loss:	1.276	Avg	Loss:	1.251
Epoch	7	Batch	1500	Loss:	0.834	Avg	Loss:	1.251
Epoch	7	Train		Loss:	1.251	Test	Loss:	1.204 Accuracy 56.75%
Epoch	8	Batch	0	Loss:	1.105	Avg	Loss:	1.105
Epoch	8	Batch	250	Loss:	1.348	Avg	Loss:	1.238
Epoch	8	Batch	500	Loss:	0.974	Avg	Loss:	1.233
Epoch	8	Batch	750	Loss:	1.127	Avg	Loss:	1.236
Epoch	8	Batch	1000	Loss:	0.960	Avg	Loss:	1.233
Epoch	8	Batch	1250	Loss:	1.234	Avg	Loss:	1.234
Epoch	8	Batch	1500	Loss:	1.000	Avg	Loss:	1.235
Epoch	8	Train		Loss:	1.235	Test	Loss:	1.222 Accuracy 57.17%
Epoch	9	Batch	0	Loss:	0.954	Avg	Loss:	0.954
Epoch	9	Batch	250	Loss:	0.920	Avg	Loss:	1.207
Epoch	9	Batch	500	Loss:	1.259	Avg	Loss:	1.225
Epoch	9	Batch	750	Loss:	1.603	Avg	Loss:	1.225
Epoch	9	Batch	1000	Loss:	0.929	Avg	Loss:	1.226
Epoch	9	Batch	1250	Loss:	0.866	Avg	Loss:	1.223
Epoch	9	Batch	1500	Loss:	1.212	Avg	Loss:	1.221
Epoch	9	Train		Loss:	1.221	Test	Loss:	1.149 Accuracy 59.18%
Epoch	10	Batch	0	Loss:	1.056	Avg	Loss:	1.056
Epoch	10	Batch	250	Loss:	1.271	Avg	Loss:	1.213
Epoch	10	Batch	500	Loss:	1.319	Avg	Loss:	1.203
Epoch	10	Batch	750	Loss:	1.218	Avg	Loss:	1.203
Epoch	10	Batch	1000	Loss:	1.276	Avg	Loss:	1.203
Epoch	10	Batch	1250	Loss:	1.275	Avg	Loss:	1.205
Epoch	10	Batch	1500	Loss:	1.021	Avg	Loss:	1.204
Epoch	10	Train		Loss:	1.205	Test	Loss:	1.176 Accuracy 58.60%
Epoch	11	Batch	0	Loss:	1.598	Avg	Loss:	1.598
Epoch	11	Batch	250	Loss:	1.187	Avg	Loss:	1.193
Epoch	11	Batch	500	Loss:	1.723	Avg	Loss:	1.202
Epoch	11	Batch	750	Loss:	1.348	Avg	Loss:	1.193
Epoch	11	Batch	1000	Loss:	1.345	Avg	Loss:	1.192
Epoch	11	Batch	1250	Loss:	1.115	Avg	Loss:	1.192
Epoch	11	Batch	1500	Loss:	1.558	Avg	Loss:	1.194
Epoch	11	Train		Loss:	1.194	Test	Loss:	1.138 Accuracy 60.02%
Epoch	12	Batch	0	Loss:	1.181	Avg	Loss:	1.181
Epoch	12	Batch	250	Loss:	1.335	Avg	Loss:	1.174
Epoch	12	Batch	500	Loss:	1.175	Avg	Loss:	1.176
Epoch	12	Batch	750	Loss:	1.431	Avg	Loss:	1.177
Epoch	12	Batch	1000	Loss:	1.098	Avg	Loss:	1.176
Epoch	12	Batch	1250	Loss:	1.361	Avg	Loss:	1.173


```

Epoch 12 Batch 1500 Loss: 1.206 Avg Loss: 1.178
Epoch 12 Train Loss: 1.179 Test Loss: 1.156 Accuracy 58.89%
Epoch 13 Batch 0 Loss: 1.334 Avg Loss: 1.334
Epoch 13 Batch 250 Loss: 1.412 Avg Loss: 1.164
Epoch 13 Batch 500 Loss: 1.574 Avg Loss: 1.165
Epoch 13 Batch 750 Loss: 1.704 Avg Loss: 1.165
Epoch 13 Batch 1000 Loss: 1.288 Avg Loss: 1.169
Epoch 13 Batch 1250 Loss: 1.082 Avg Loss: 1.171
Epoch 13 Batch 1500 Loss: 1.254 Avg Loss: 1.167
Epoch 13 Train Loss: 1.166 Test Loss: 1.146 Accuracy 59.68%
Epoch 14 Batch 0 Loss: 1.304 Avg Loss: 1.304
Epoch 14 Batch 250 Loss: 0.758 Avg Loss: 1.132
Epoch 14 Batch 500 Loss: 1.011 Avg Loss: 1.155
Epoch 14 Batch 750 Loss: 1.062 Avg Loss: 1.156
Epoch 14 Batch 1000 Loss: 1.247 Avg Loss: 1.160
Epoch 14 Batch 1250 Loss: 1.869 Avg Loss: 1.162
Epoch 14 Batch 1500 Loss: 1.162 Avg Loss: 1.166
Epoch 14 Train Loss: 1.163 Test Loss: 1.108 Accuracy 60.84%
Epoch 15 Batch 0 Loss: 1.372 Avg Loss: 1.372
Epoch 15 Batch 250 Loss: 1.278 Avg Loss: 1.158
Epoch 15 Batch 500 Loss: 0.995 Avg Loss: 1.158
Epoch 15 Batch 750 Loss: 1.159 Avg Loss: 1.154
Epoch 15 Batch 1000 Loss: 0.829 Avg Loss: 1.149
Epoch 15 Batch 1250 Loss: 0.891 Avg Loss: 1.150
Epoch 15 Batch 1500 Loss: 1.061 Avg Loss: 1.153
Epoch 15 Train Loss: 1.152 Test Loss: 1.122 Accuracy 60.60%
Epoch 16 Batch 0 Loss: 1.154 Avg Loss: 1.154
Epoch 16 Batch 250 Loss: 1.100 Avg Loss: 1.135
Epoch 16 Batch 500 Loss: 1.220 Avg Loss: 1.136
Epoch 16 Batch 750 Loss: 1.080 Avg Loss: 1.135
Epoch 16 Batch 1000 Loss: 1.006 Avg Loss: 1.140
Epoch 16 Batch 1250 Loss: 1.313 Avg Loss: 1.140
Epoch 16 Batch 1500 Loss: 0.844 Avg Loss: 1.141
Epoch 16 Train Loss: 1.142 Test Loss: 1.114 Accuracy 60.65%
Epoch 17 Batch 0 Loss: 1.206 Avg Loss: 1.206
Epoch 17 Batch 250 Loss: 0.983 Avg Loss: 1.127
Epoch 17 Batch 500 Loss: 1.257 Avg Loss: 1.134
Epoch 17 Batch 750 Loss: 0.889 Avg Loss: 1.136
Epoch 17 Batch 1000 Loss: 1.154 Avg Loss: 1.137
Epoch 17 Batch 1250 Loss: 1.087 Avg Loss: 1.142
Epoch 17 Batch 1500 Loss: 1.229 Avg Loss: 1.139
Epoch 17 Train Loss: 1.138 Test Loss: 1.107 Accuracy 60.72%
Epoch 18 Batch 0 Loss: 0.840 Avg Loss: 0.840
Epoch 18 Batch 250 Loss: 1.138 Avg Loss: 1.112
Epoch 18 Batch 500 Loss: 1.337 Avg Loss: 1.119
Epoch 18 Batch 750 Loss: 1.284 Avg Loss: 1.125
Epoch 18 Batch 1000 Loss: 0.942 Avg Loss: 1.117
Epoch 18 Batch 1250 Loss: 0.880 Avg Loss: 1.119
Epoch 18 Batch 1500 Loss: 1.076 Avg Loss: 1.120
Epoch 18 Train Loss: 1.119 Test Loss: 1.095 Accuracy 60.96%
Epoch 19 Batch 0 Loss: 1.066 Avg Loss: 1.066
Epoch 19 Batch 250 Loss: 1.449 Avg Loss: 1.090
Epoch 19 Batch 500 Loss: 1.107 Avg Loss: 1.113
Epoch 19 Batch 750 Loss: 1.576 Avg Loss: 1.119
Epoch 19 Batch 1000 Loss: 1.165 Avg Loss: 1.122
Epoch 19 Batch 1250 Loss: 0.922 Avg Loss: 1.122
Epoch 19 Batch 1500 Loss: 0.991 Avg Loss: 1.120
Epoch 19 Train Loss: 1.120 Test Loss: 1.096 Accuracy 61.58%


```

Plot the training loss (and validation loss/accuracy, if recorded).

```

In [28]: plt.plot(per_epoch_test_loss)
plt.plot(per_epoch_train_loss)
plt.plot([x/100. for x in per_epoch_accuracy])
plt.legend(["Test Loss", "Train Loss", "Accuracy"])
plt.show()

```



In []: